LINEAR MODELS ARE GREAT

ALEX GOLD
SOLUTIONS ENGINEER
RSTUDIO

 @alexkgold

Slides at: https://github.com/akgold/dsdc_linear_models

IT ME.

WHAT DOES IT MEAN TO BE LINEAR?

21.0
21.0
22.8
21.4
18.7
18.1
14.3
24.4
22.8
19.2
17.8
16.4
17.3
15.2
10.4
10.4
14.7
32.4
30.4
33.9
21.5
15.5
15.2
13.3
19.2
27.3
26.0
30.4
15.8
19.7
15.0
21.4

$$= \beta_0 + \beta_{cyl}$$

Just a number

6
6
4
6
8
6
8
4
4
6
6
8
8
8
8
8
8
4
4
4
4
8
8
8
8
4
4
4
8
6
8
4

$$+ \beta_{hp}$$

160.0
160.0
108.0
258.0
360.0
225.0
360.0
146.7
140.8
167.6
167.6
275.8
275.8
275.8
472.0
460.0
440.0
78.7
75.7
71.1
120.1
318.0
304.0
350.0
400.0
79.0
120.3
95.1
351.0
145.0
301.0
121.0

+ some prediction error

WHAT DOES IT MEAN TO BE LINEAR?

$$Y = \beta X + \epsilon$$

ALEX <3 LINEAR MODELS

```
> lm(mpg ~ cyl + disp + hp, data = mtcars)

Call:
lm(formula = mpg ~ cyl + disp + hp, data = mtcars)

Coefficients:
(Intercept)          cyl         disp           hp
   34.18492     -1.22742     -0.01884     -0.01468
```

$$Y = \beta X + \epsilon$$

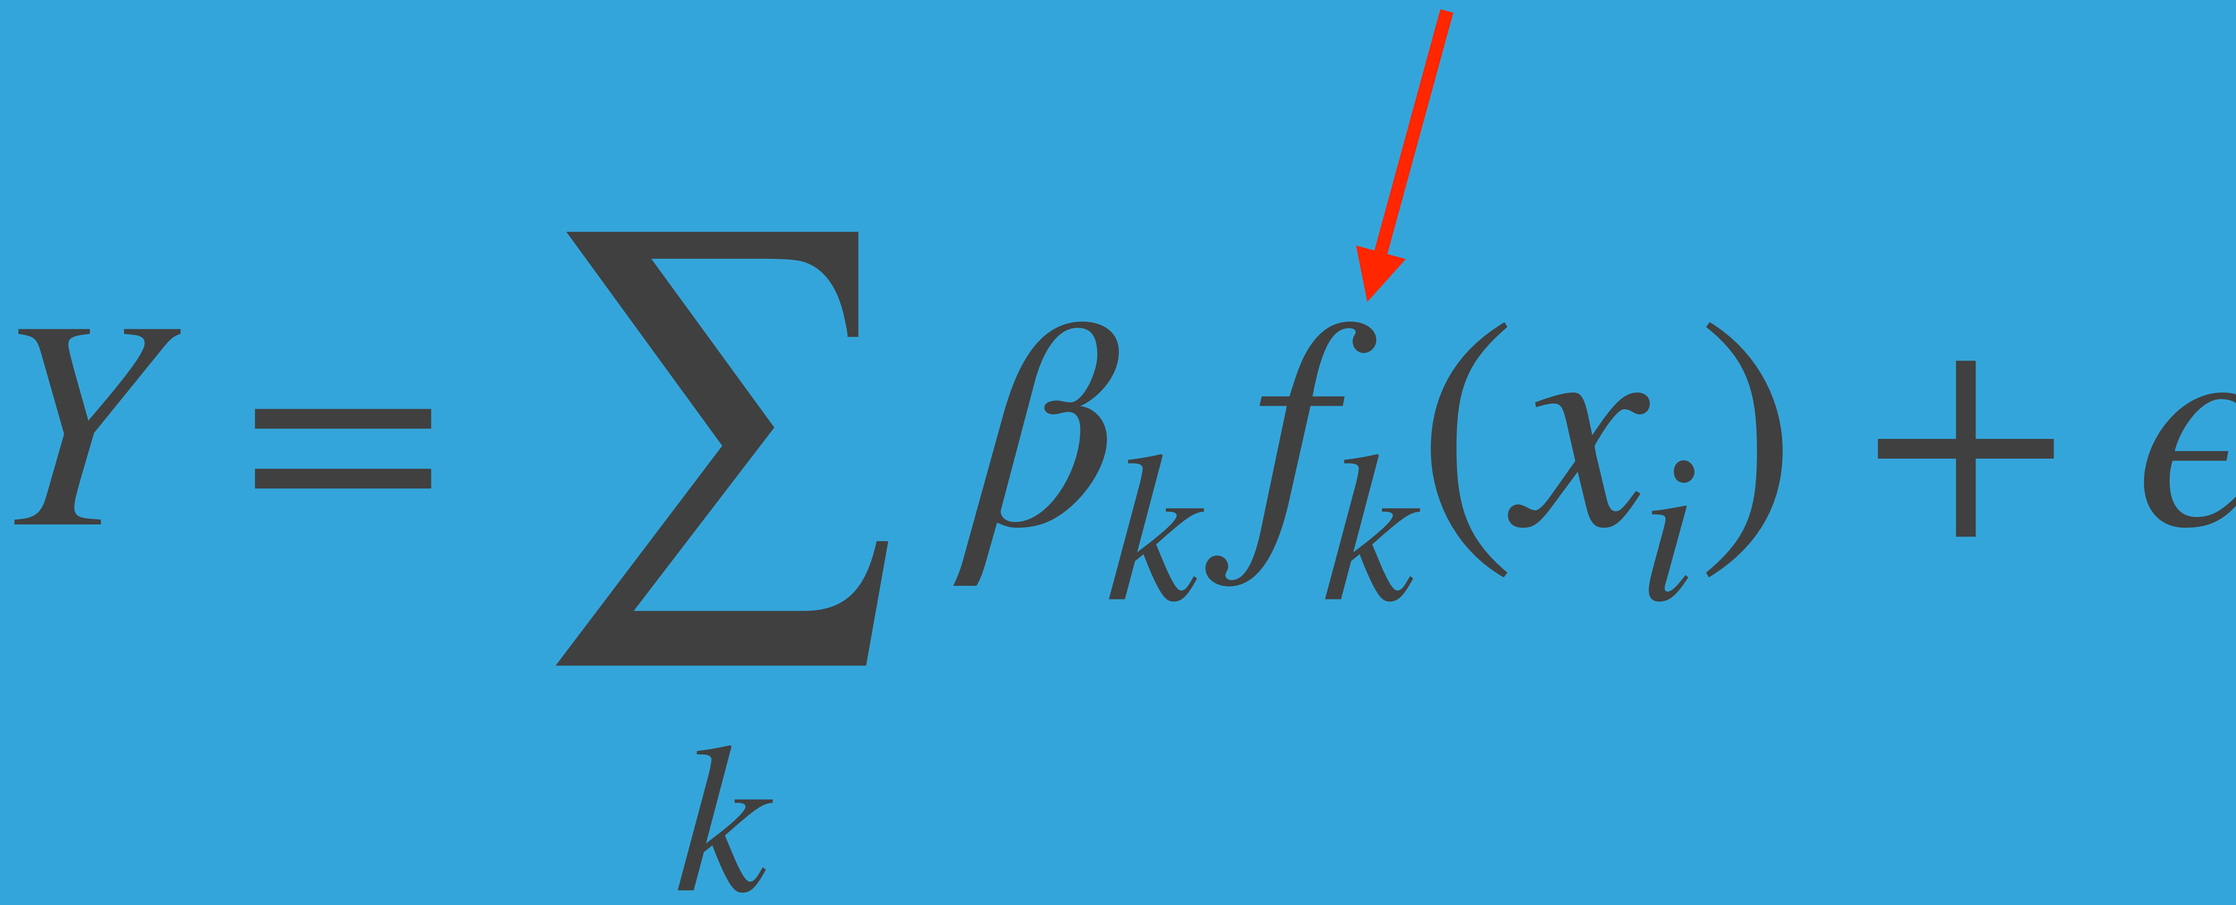
$Y = \sum_k \beta_k f_k(x_i) + \epsilon$

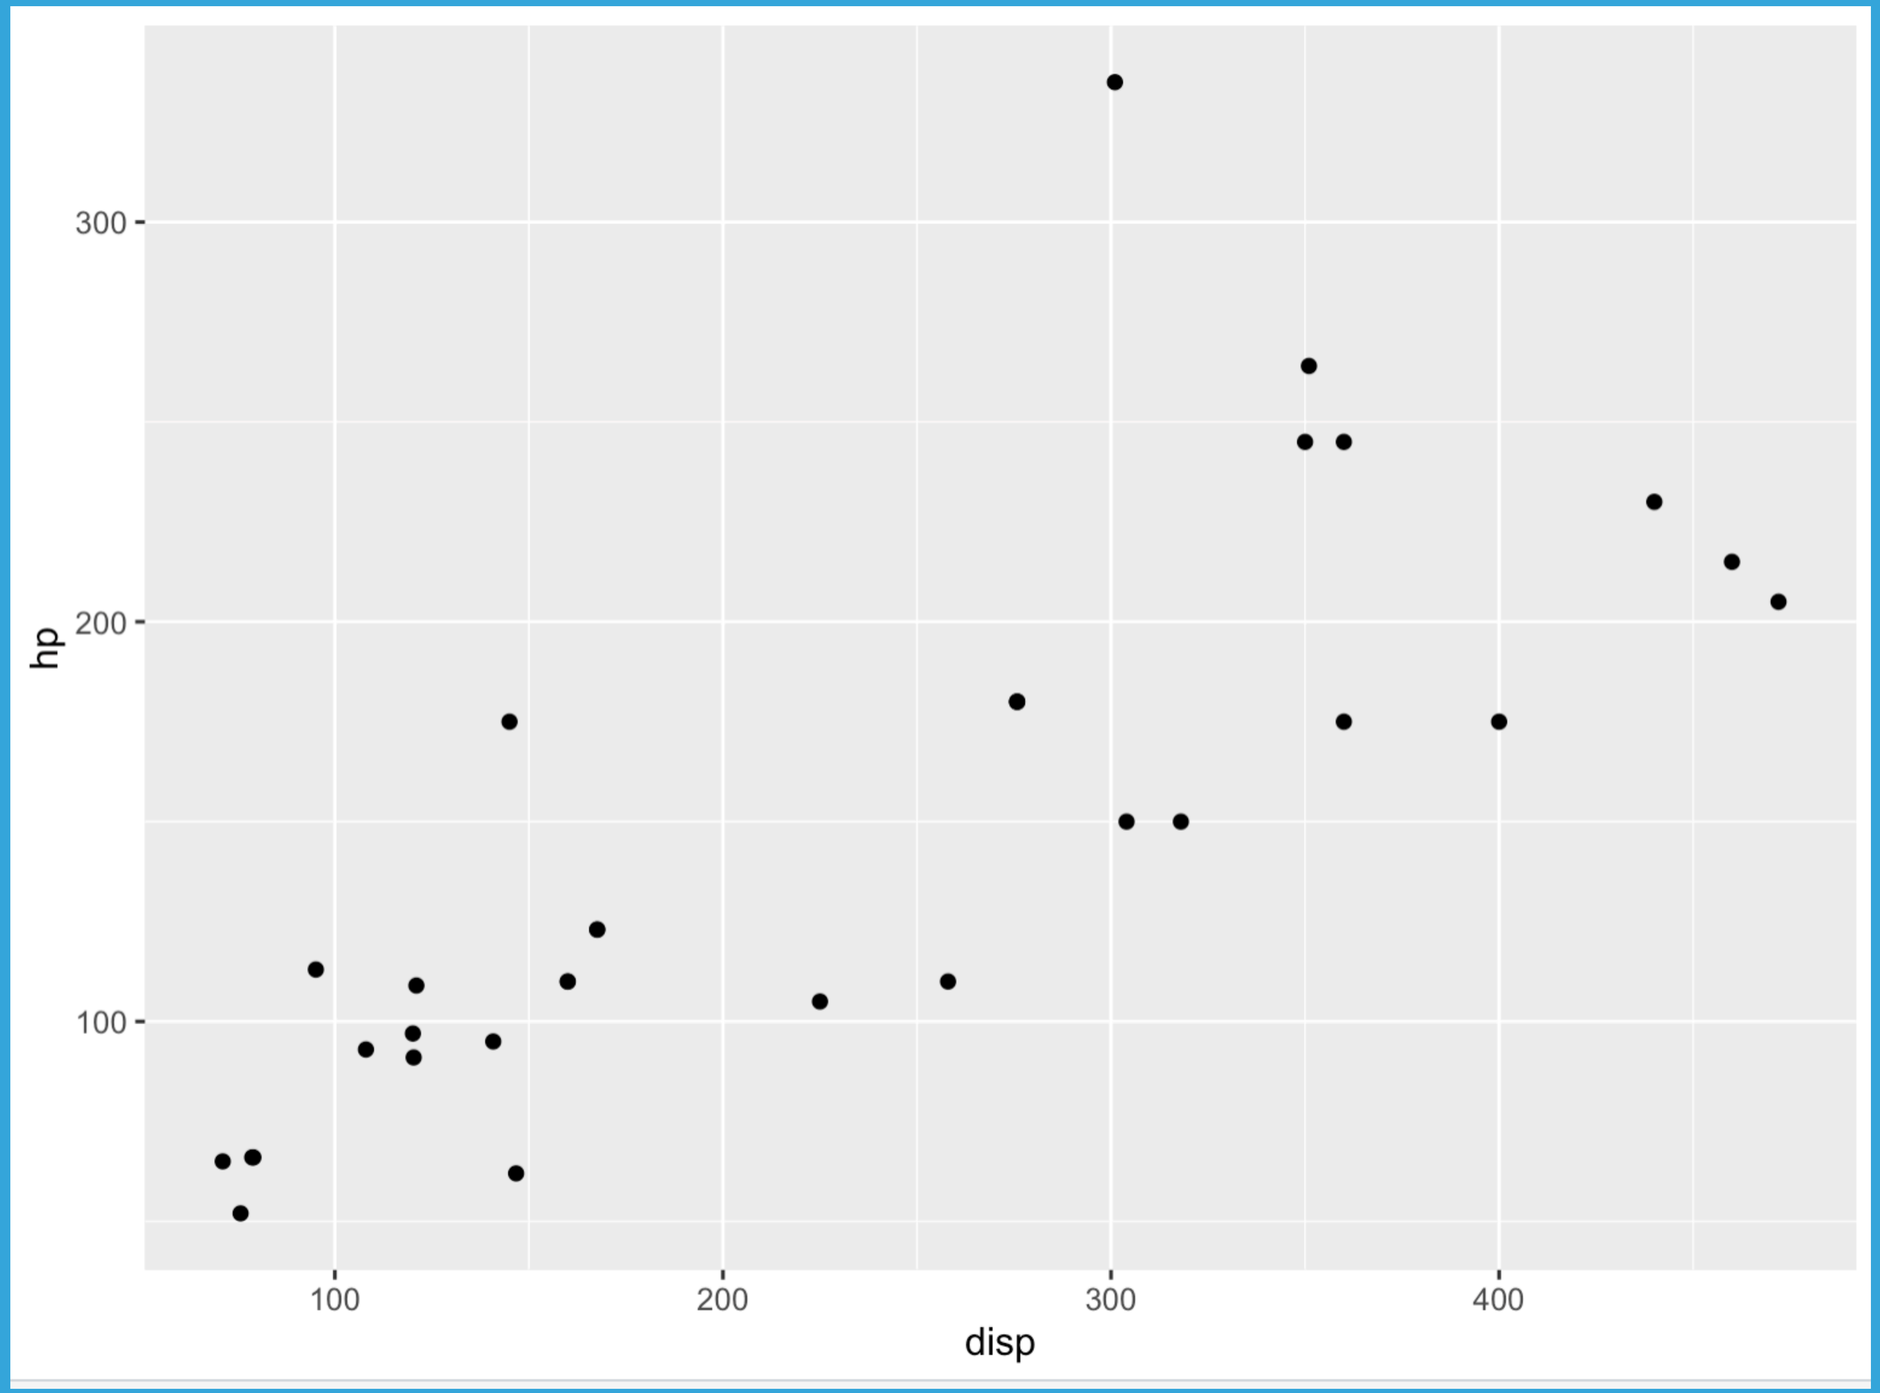
300
200
100
hp
100
200
300
400
disp

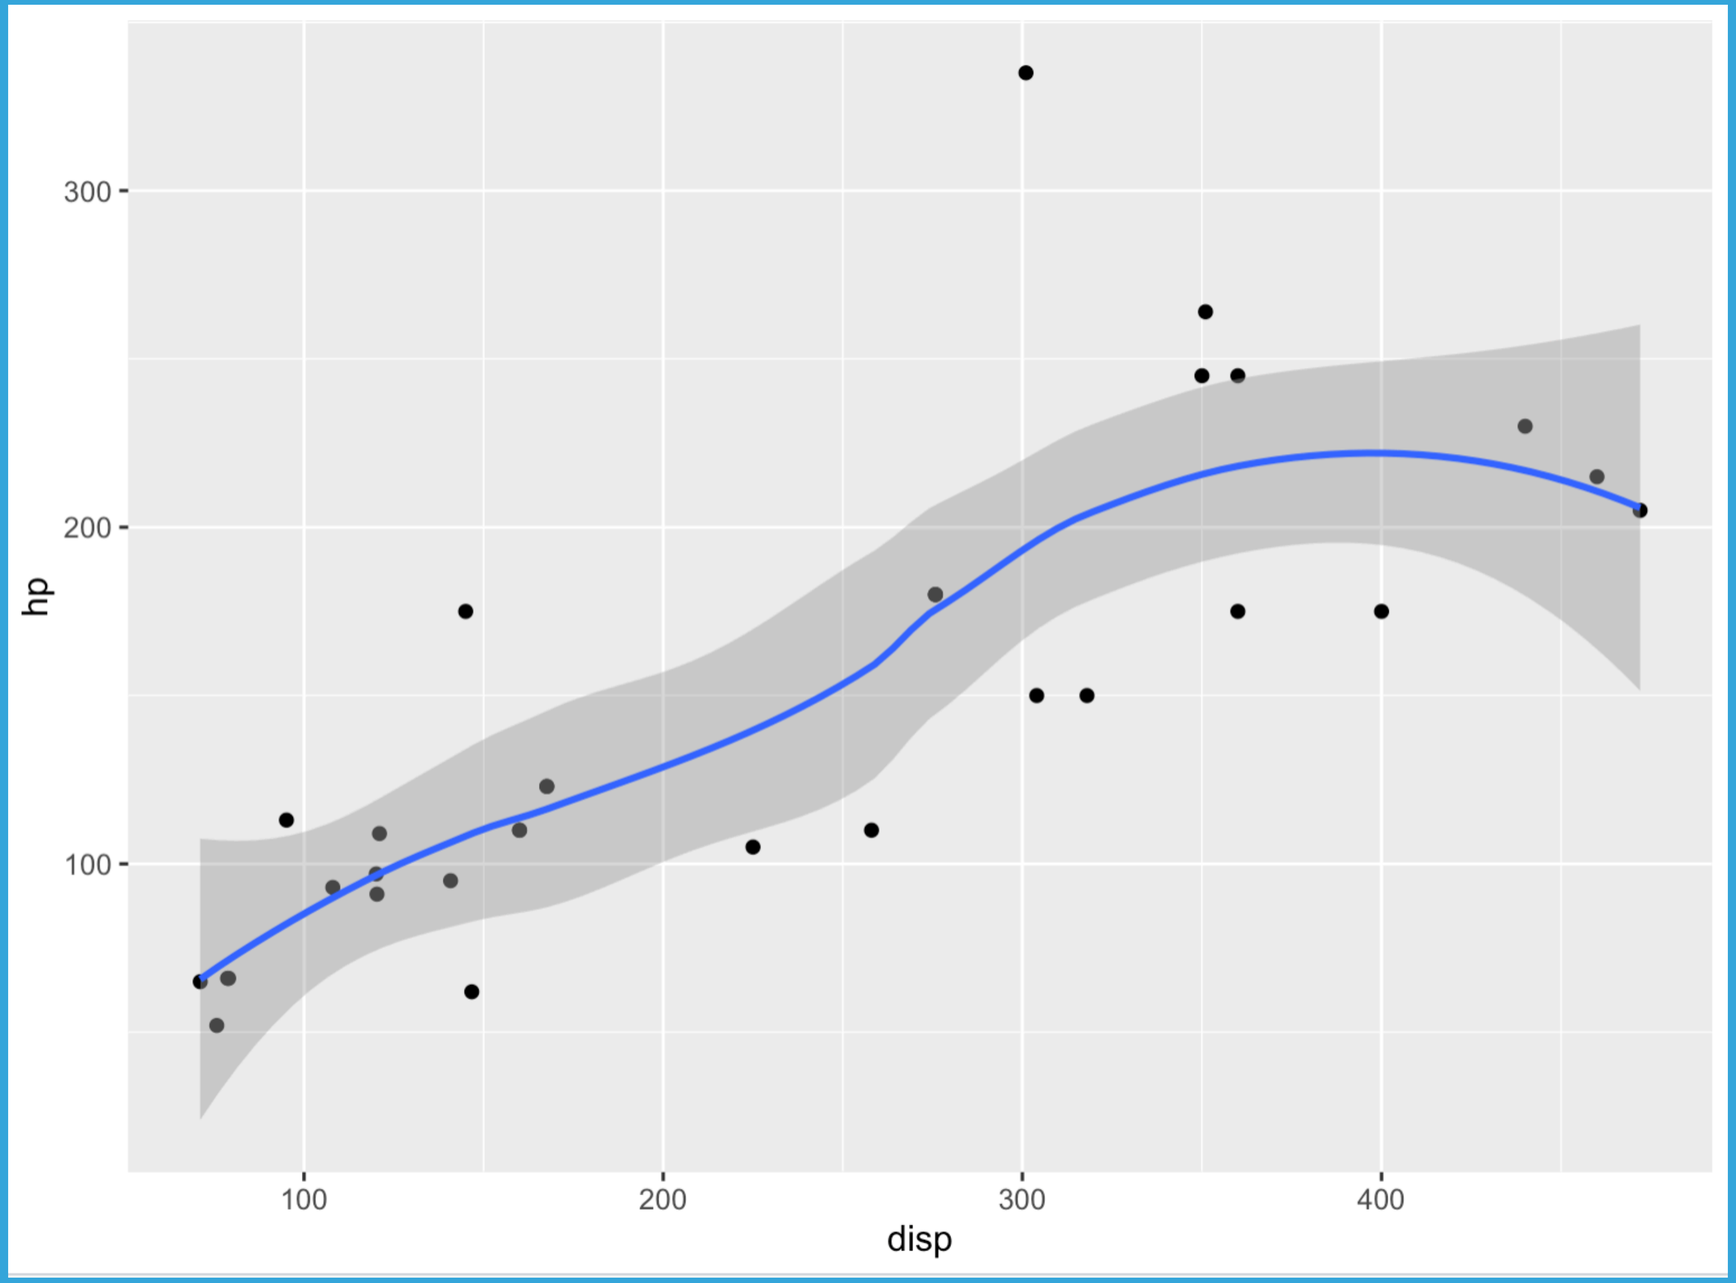


```
ggplot(mtcars, aes(x = disp, y = hp)) + geom_point() + stat_smooth(method = "loess")
```

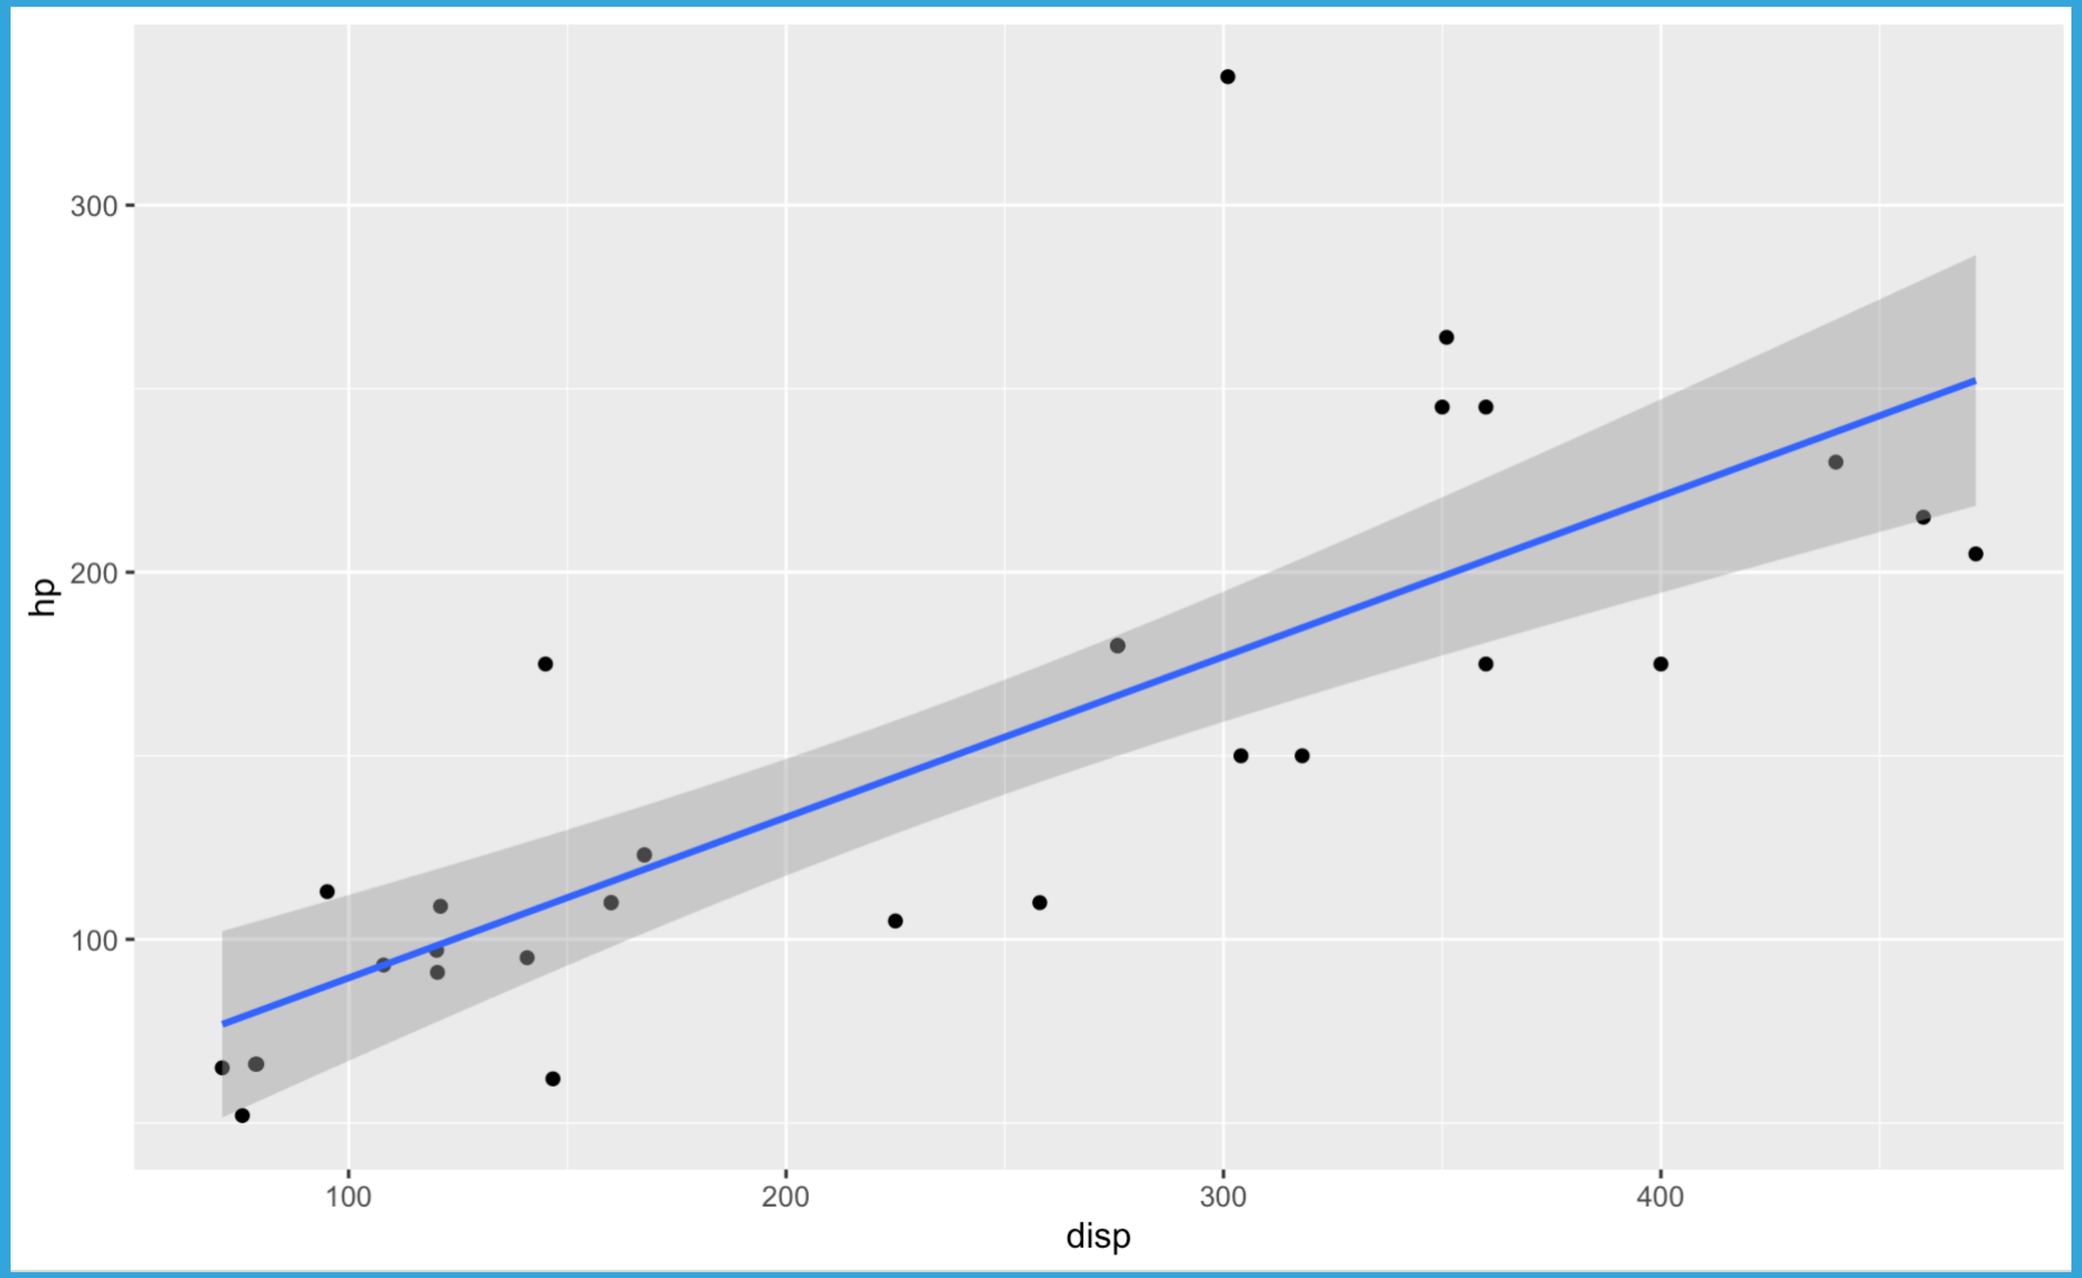


```
ggplot(mtcars, aes(x = disp, y = hp)) + geom_point() + stat_smooth(method = "lm")
```

$$gender = \beta_0 + \beta_1 weight + \beta_2 height$$

$$gender = \beta_0 + \beta_1 weight + \beta_2 height$$

vs

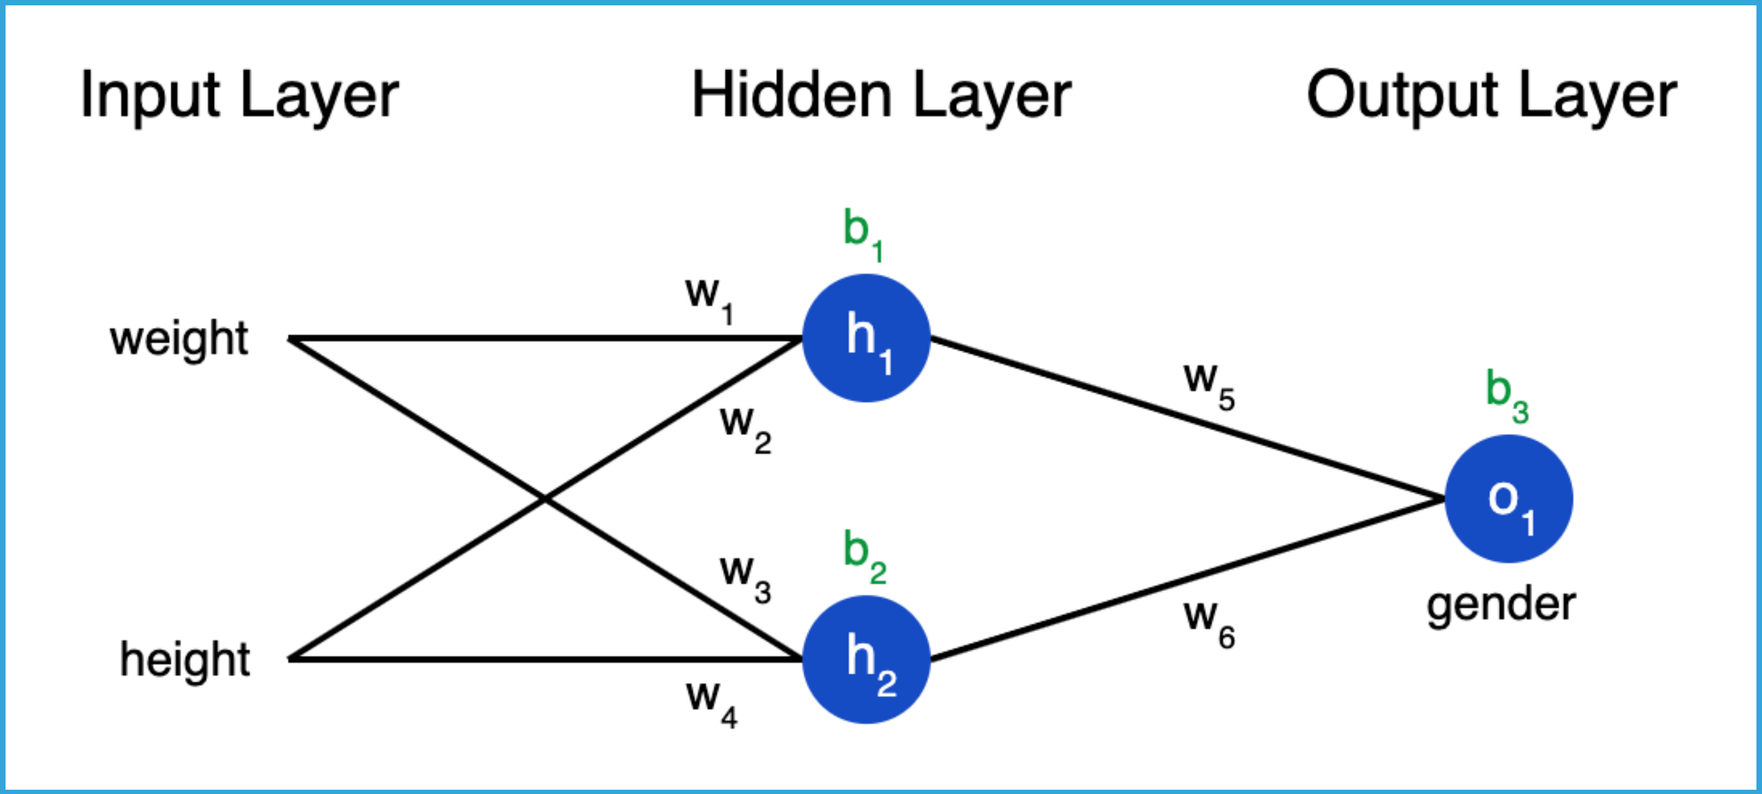

$$gender = \beta_0 + \beta_1 weight + \beta_2 height$$

vs

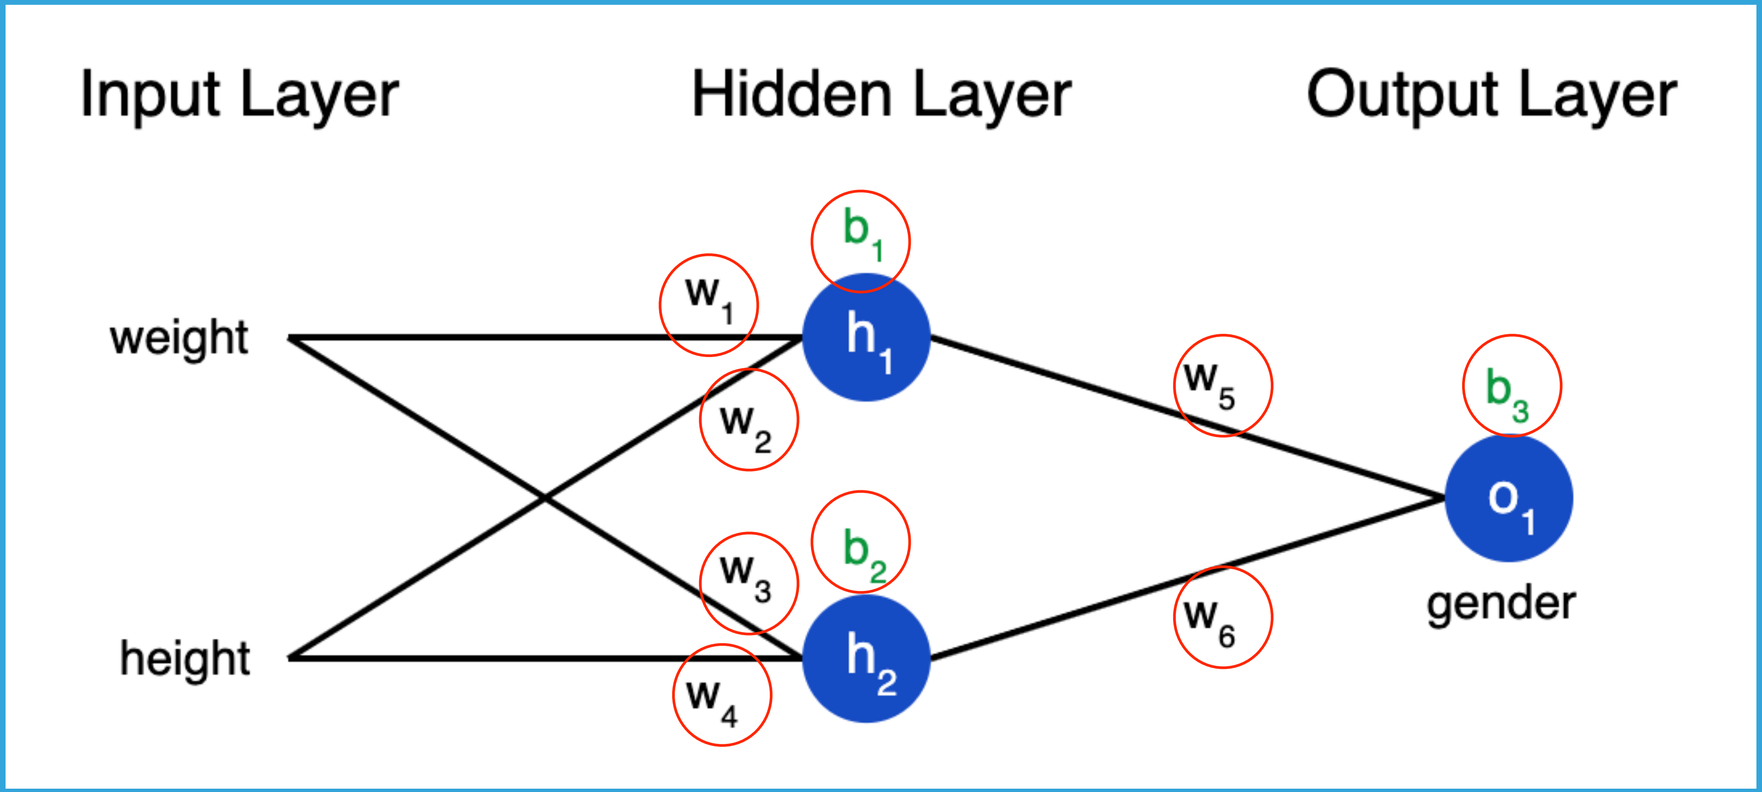

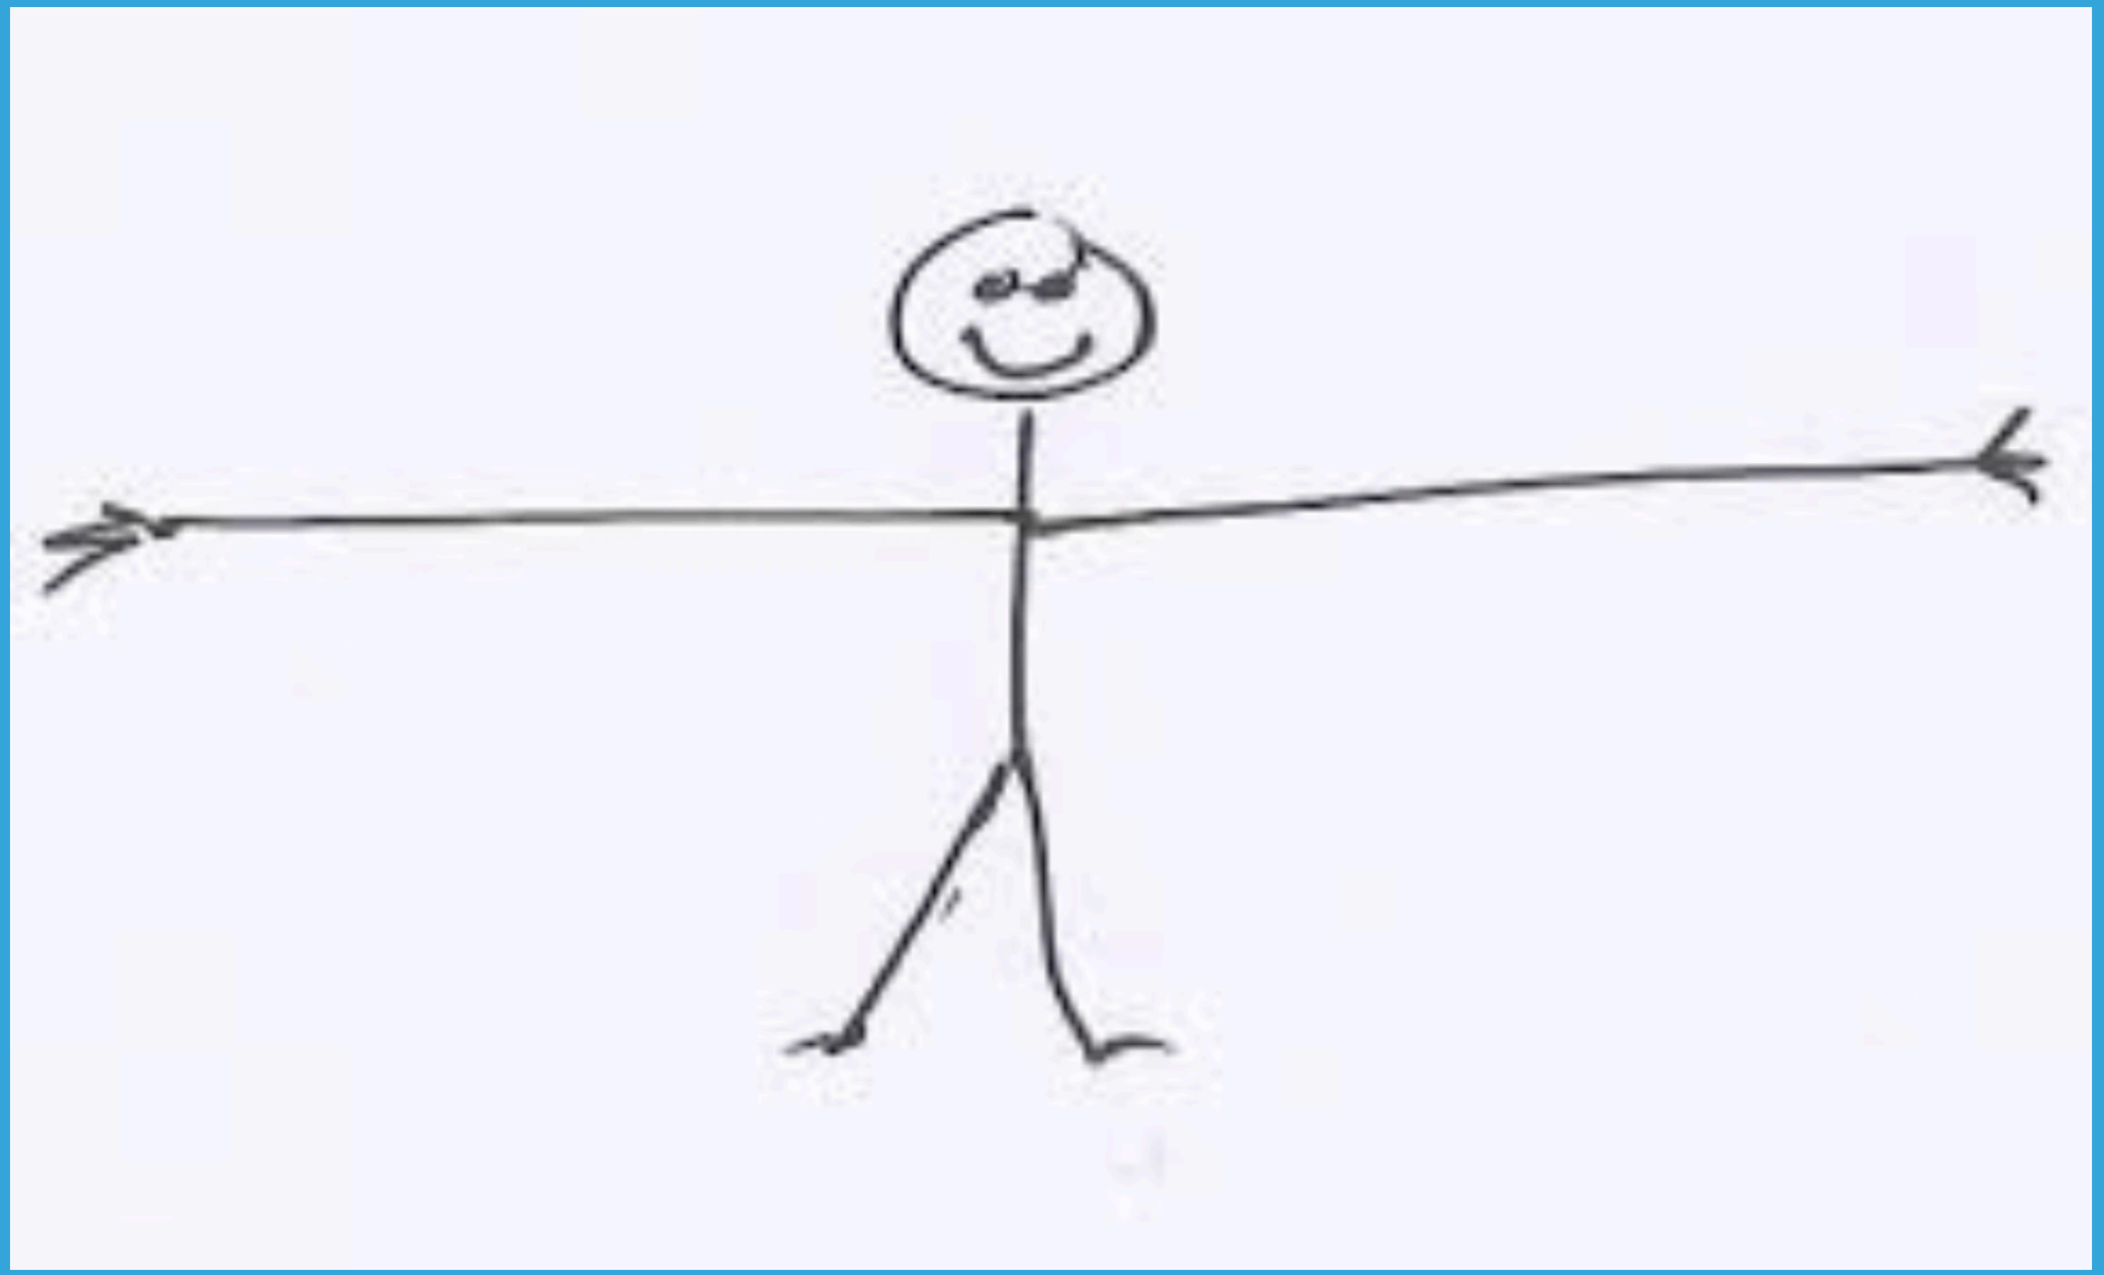

It’s all about the Data-Generating Process

$$mpg = \beta_0 + \beta_1 cyl + disp$$

OR

$$mpg = \beta_0 + \beta_1 cyl + \beta_2 cyl^2 + \beta_3 log(disp)$$

?

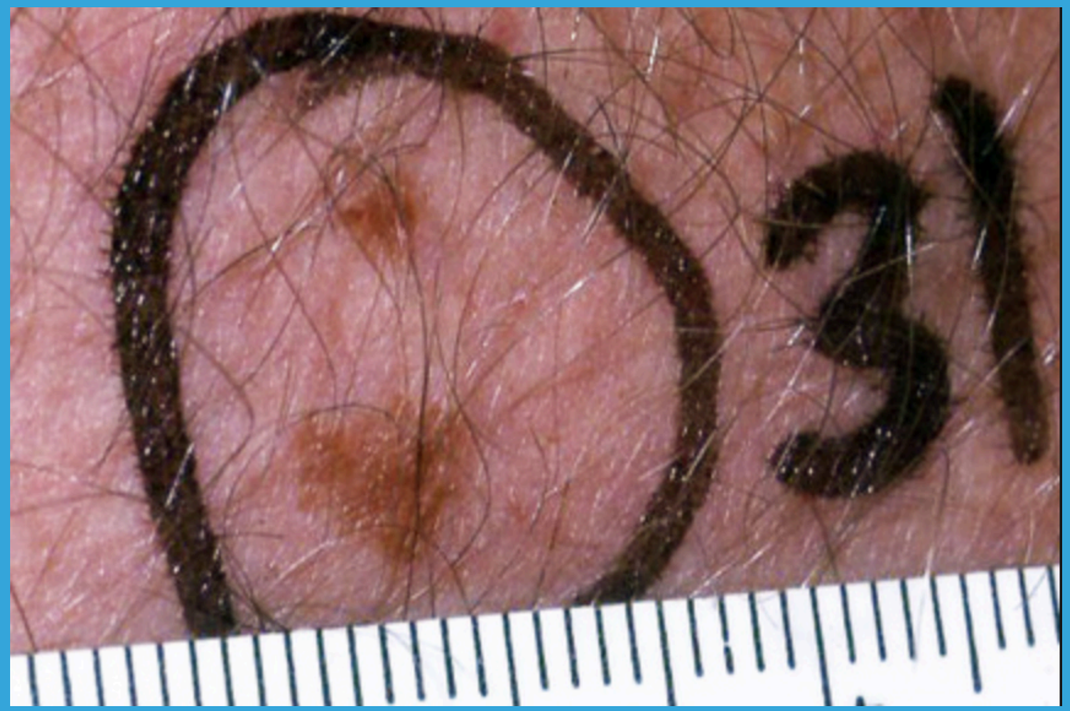
031

1. INTERPRETATION MATTERS.
2. LINEARITY ISN'T RESTRICTIVE.
3. MO' SQUIGGLY = MO' OVERFITTING.
4. SMALL DATA'S OK.
5. IT'S ALL ABOUT THE DGP.

@alexkgold

https://github.com/akgold/dsdc_linear_models